Lecture 20 - Mar 18

Clustering
Expectation-Maximization (EM) Algorithm

References

Data Mining and Machine Learning
Ch 13 - Representation-based Clustering

Elements of Statistical Learning
14.3.6 K-means
14.3.7 Gaussian Mixtures as Soft K-Means Clustering

AKA data segmentation has a variety of goals.

All group or segment objects into "clusters" where points inside are similar but different from points in other clusters.

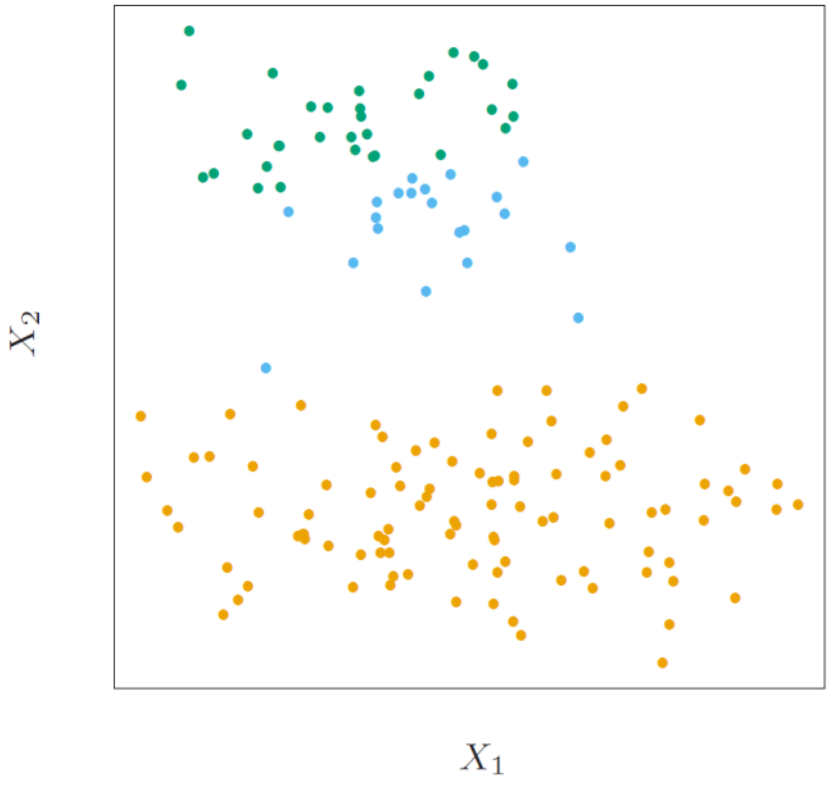

- Sometimes the goal is to find a natural hierarchy of clusters

- Also used for descriptive statistics regarding whether or not there exist distinct subgroups with substantially different properties

Clustering Algorithms

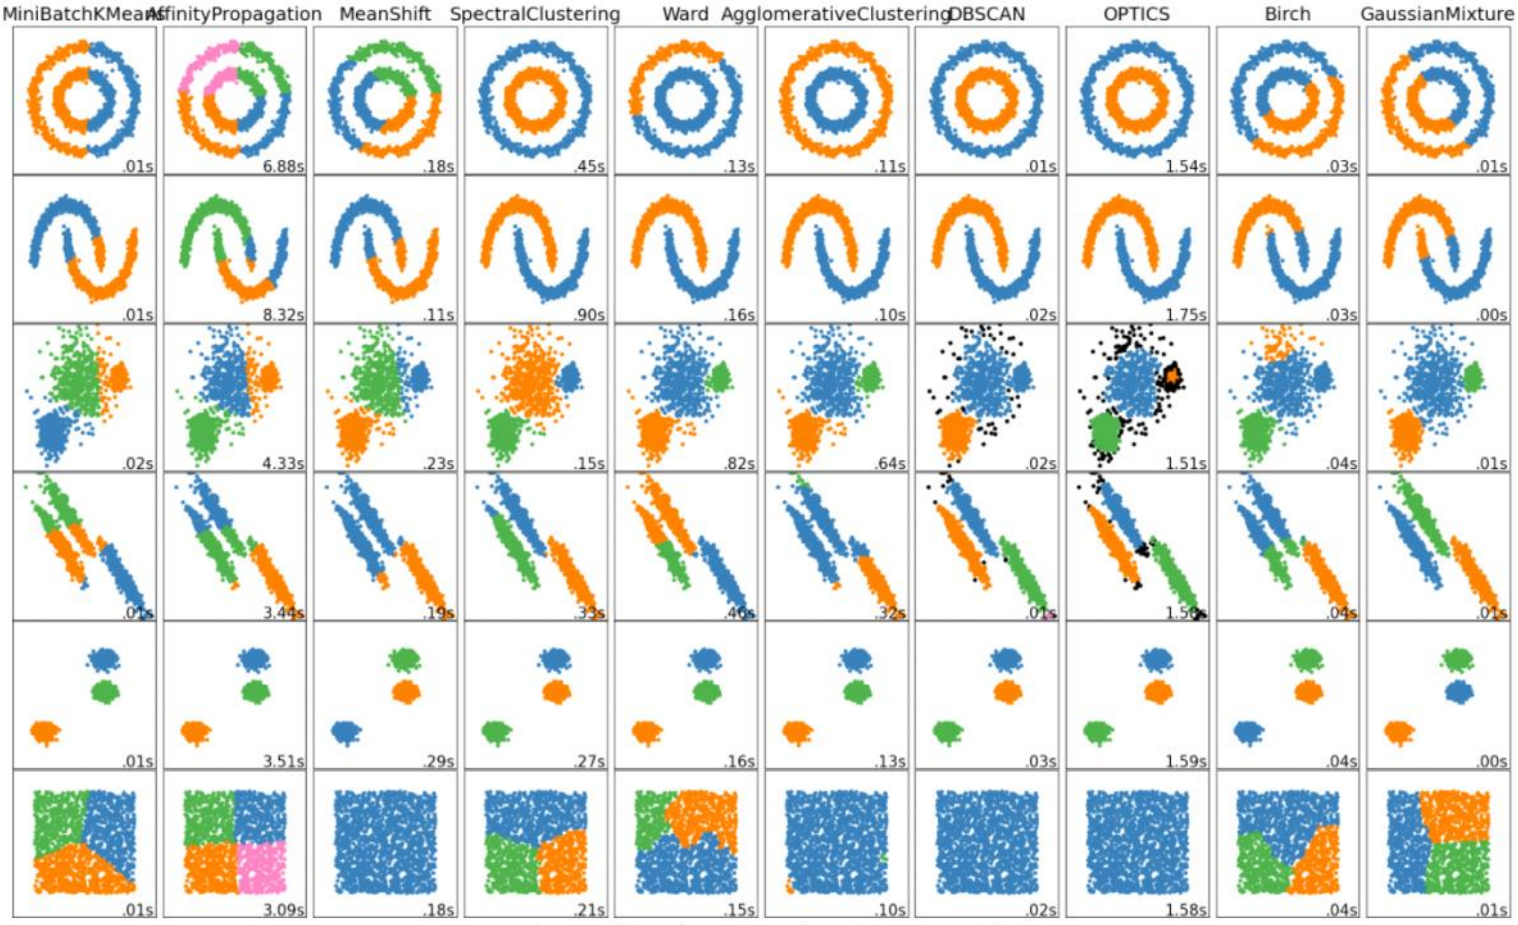

A comparison of the clustering algorithms in scikit-learn

https://scikit-learn.org/stable/modules/clustering.html

Types: Combinatorial, Mixture modeling, mode seeking

Combinatorial: work directly on the observed data

Mixture modeling: assumes data is an i.i.d. sample from a population described by a pdf characterized by a parameterized model taken as a mixture of component density functions: each describing one cluster

Mode seeking: "bump hunters" attempt to directly estimate distinct modes of the pdf as centers of clusters

Let $X \in \mathbb{R}^{n \times d}$ be a dataset of points $x_i \in \mathbb{R}^d$

Representative-based clustering partitions X into k clusters...

$$C = \{C_1, \ldots, C_k\}$$

For each cluster C_i, there exists a representative point that summarizes the cluster

↳ e.g. the centroid (mean) μ_i of points in C_i

$$\mu_i = \frac{1}{n_i} \sum_{x_j \in C_i} x_j \quad \text{where } n_i = |C_i|$$

Two common methods: K-means + Expectation-Maximization (EM) algorithms

For a given clustering C, we need a scoring function to evaluate it

The SSE scoring function is

$$SSE(C) = \sum_{i=1}^{k} \sum_{x_j \in C_i} \|x_j - m_i\|^2$$

Sum over all clusters

SS of difference between cluster pts + cluster centroid

$$C^* = \arg\min_{C} \{SSE(C)\}$$

K-means minimizes SSE via a greedy iterative approach (risks reaching local min)

K-means initializes centroids randomly + iterates 2 steps

1. Assign points to nearest centroid's cluster
2. Compute centroid as mean of points in each cluster

└ (2) Compute centroid as mean of points in cluster

Stop when centroids stop moving.

Often, people run K-means several times due to random initializations + keep the best result (smallest SSE)

K-means creates convex-shaped clusters.

Algorithm 13.1

Input: X, k, ε

often, sample uniformly in each feature

1. Randomly initialize k centroids $\mu_1^0, \ldots, \mu_k^0 \in \mathbb{R}^d$

while $\sum_{i=1}^{k} \|\mu_i^t - \mu_i^{t-1}\| > \varepsilon$:

while centroids move substantially in re-calculations

- for $x_j \in X$:
  - $i^* = \arg\min_i \{\|x_j - \mu_i^{t-1}\|^2\}$
  - $C_{i^*} = C_{i^*} \cup \{x_j\}$

assign points to cluster with nearest centroid

- for $C_i \in C$
  - $\mu_i^t = \frac{1}{n_i} \sum_{x_j \in C_i} x_j$

re-calculate centroids with new clusters

K-Means in 1D

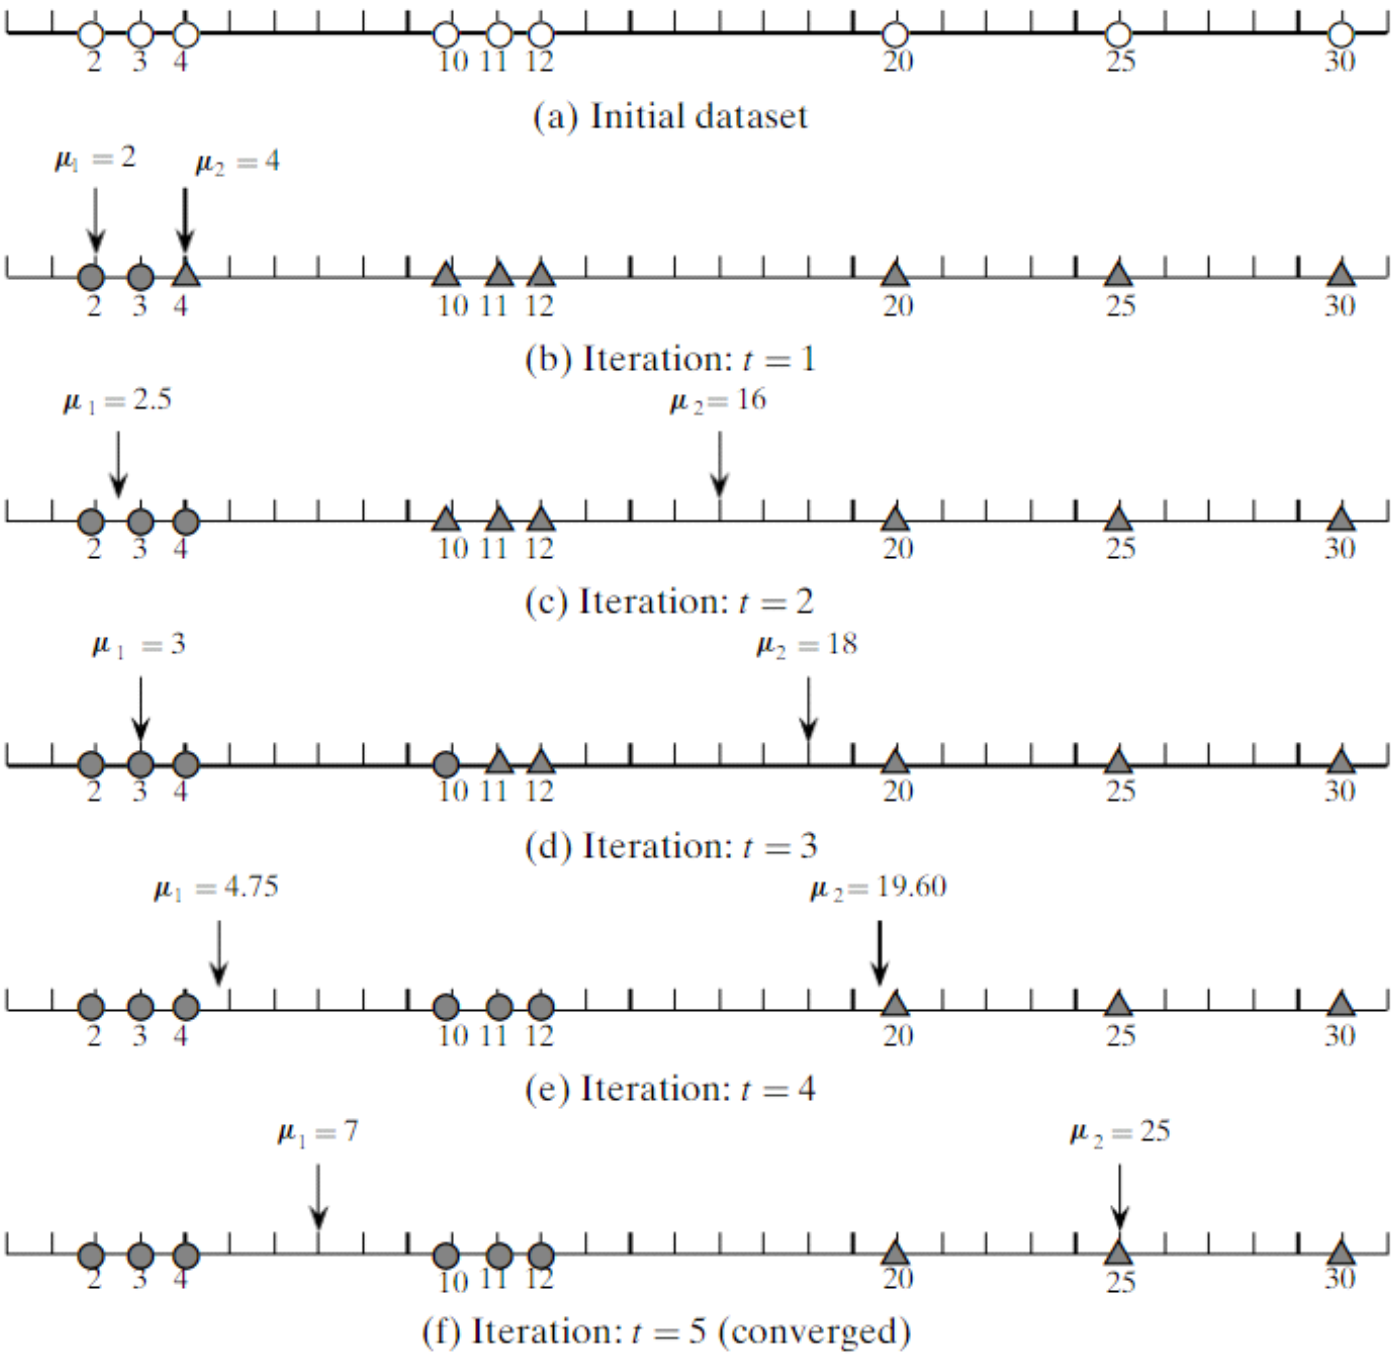

Figure 13.1. K-means in one dimension.

K-Means in 2D

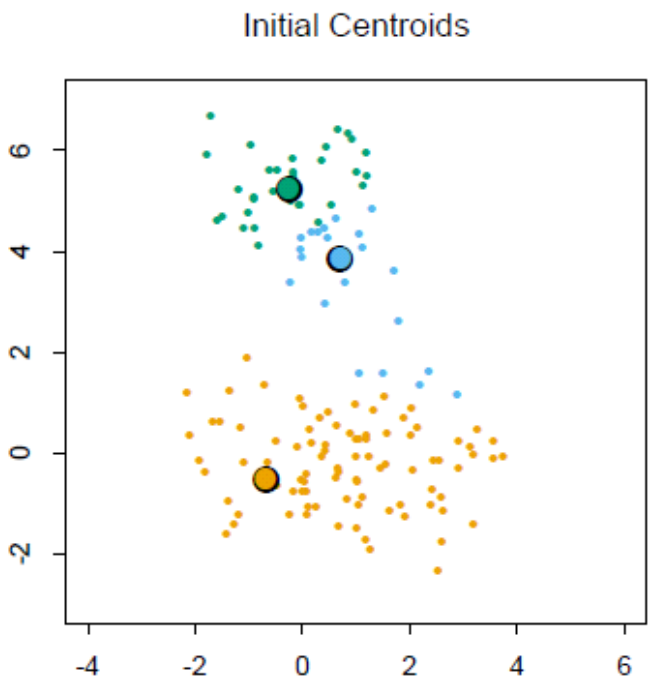
Initial Centroids
6
4
2
0
-2
-4
-2
0
2
4
6

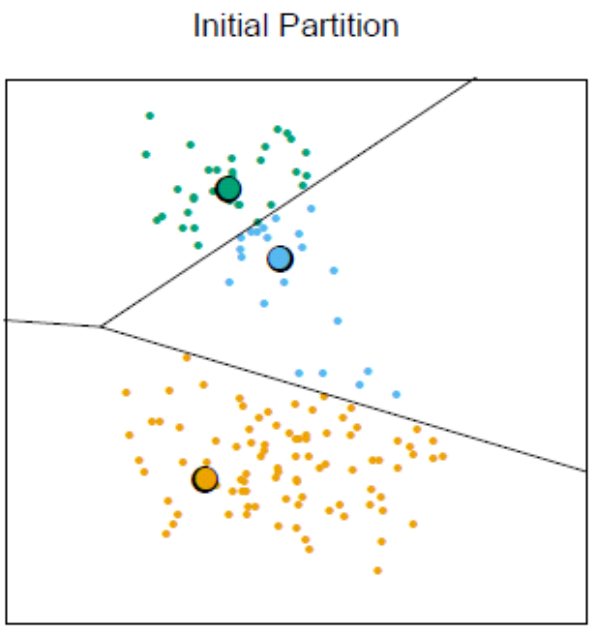
Initial Partition

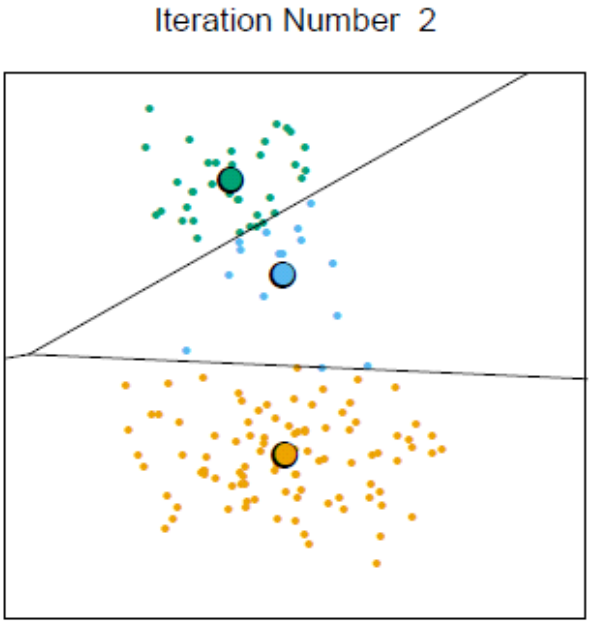
Iteration Number 2

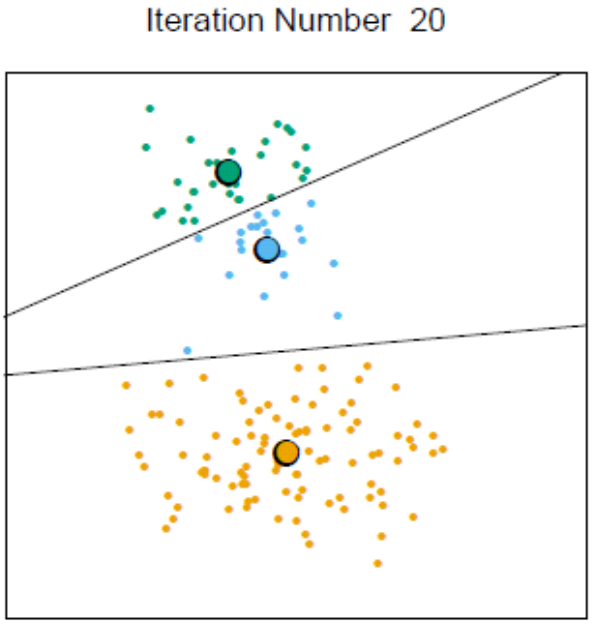
Iteration Number 20